

DevTools GTC 2020 Lab Content

Requirements	2
Story	2
Recordings	3
Setup	4
Host Setup	4
Step 0: Nsight Eclipse Edition (plugin)	5
Step 1: cuda-gdb	8
Step 2: Compute Sanitizer*	11
Steps 3-5: Nsight Systems	14
Step 3: Analyze original application behavior	14
Step 4: Buffer reuse to minimize cudaMalloc/cudaFree	16
Step 5: Rendering on Demand	18
Steps 6-7: Nsight Compute	18
Step 6.1: Analyze ExtractFeatures kernel	19
Step 6.2: Analyze the 32bit float math ExtractFeatures kernel	23
Step 7: Analyze the three optimized kernels	26

Requirements

- Ubuntu 18.04 host (other OS variants may provide different results)
- Eclipse 4.9 (2018-09) for C++ developers
- default-jre apt package
- python apt package
- CUDA 11.0 host (content compatible with CUDA 10.2)
- Turing GPU (tested on RTX 2080Ti)
 - If the device has multiple GPUs, ensure that compute (CUDA) and display (OpenGL) are on the same device, i.e. set `CUDA_VISIBLE_DEVICES` to the one with the display. Otherwise, CUDA plus OpenGL will result in high memory mapping overhead.
- libglew-dev apt package
- freeglut3-dev package

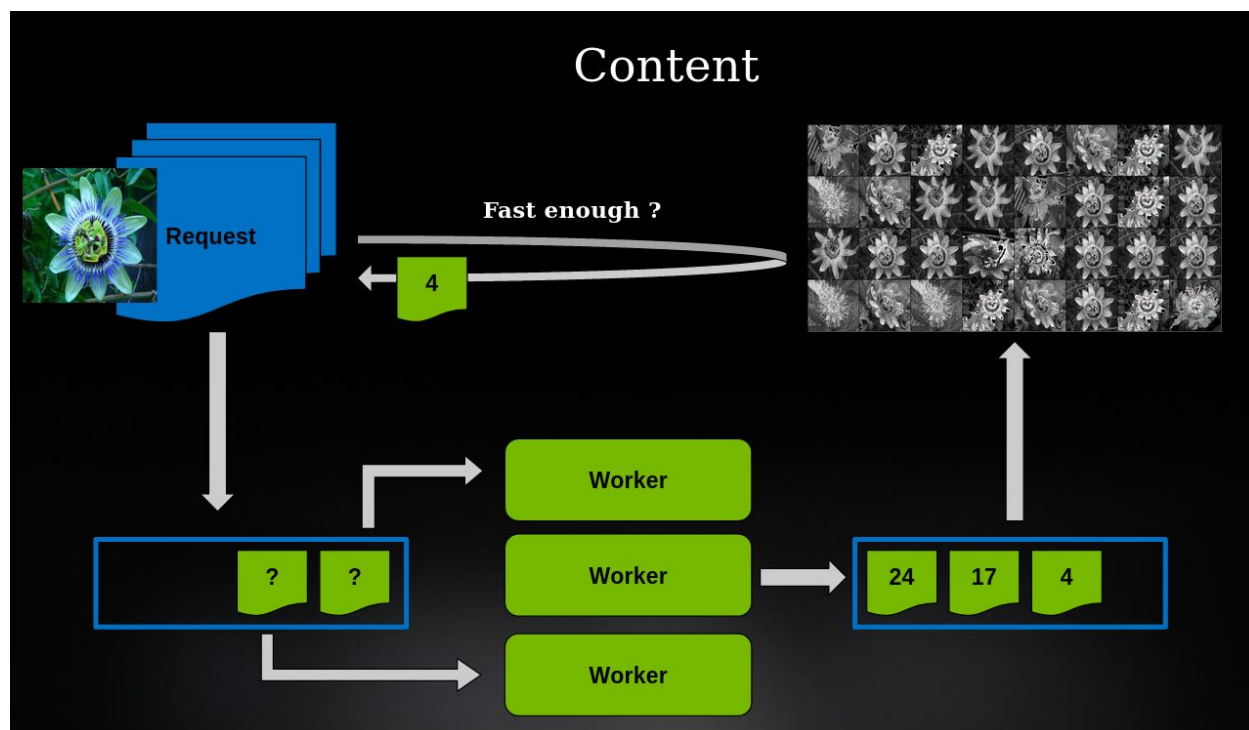
Story

The exercise represents a cloud-based image classification application. It contains a database (DB) of flower images. Users can submit their own flower pictures and the app classifies them against the DB and sends the result back to the user. The application allows users to tweak several factors including the threshold, number of source threads generating requests, and the number of worker threads processing these requests.

For simplicity, the DB contains a fixed number of flower images, and users only submit pictures of one of those flowers. For further simplification, the user images are identical to the DB versions, except for added white noise to simulate some random errors when taking the snapshot.

The exercise is to first fix several bugs in the application using correctness analysis tools to make the classification work properly. Then, users must optimize the application (framework and kernels) so that requests are served within some defined time threshold (e.g. 100ms). It is structured into steps that build on each other, but it is possible to start at any chosen step, as each one will contain all previous exercise content (e.g. all bugs are fixed in the first performance analysis exercise).

This exercise uses NVIDIA® Nsight™ Eclipse Edition but all steps can be completed when compiling and running directly from the command line. These steps are marked as “CLI Alternative” below.



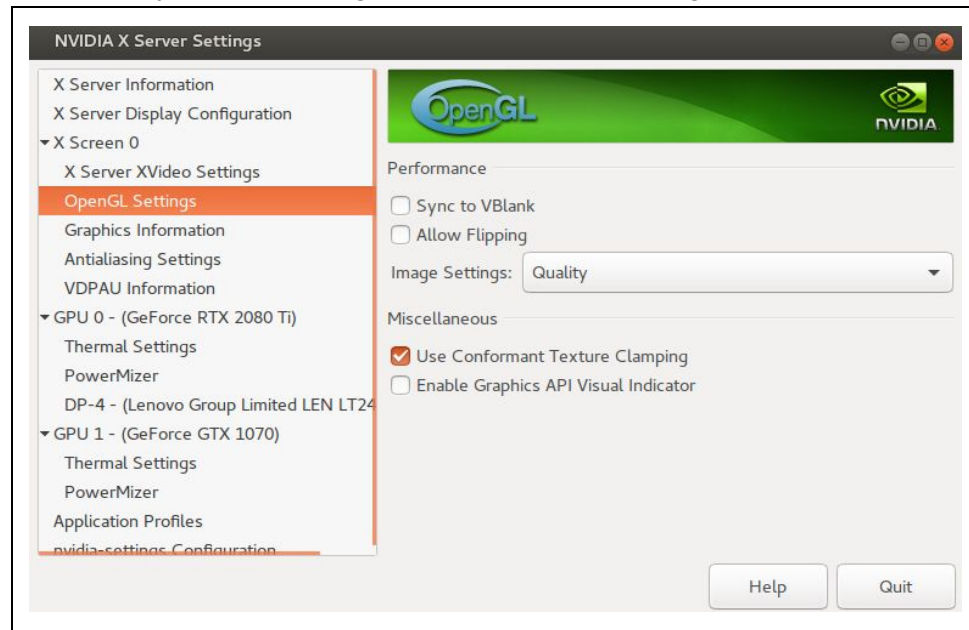
Recordings

- Step 0: Nsight Eclipse Edition <https://developer.nvidia.com/gtc/2020/video/t21395-1>
- Step 1: cuda-gdb <https://developer.nvidia.com/gtc/2020/video/t21395-2>
- Step 3-4: Nsight Systems <https://developer.nvidia.com/gtc/2020/video/t21395-3>
- Step 5-7: Nsight Compute <https://developer.nvidia.com/gtc/2020/video/t21395-4>

Setup

Host Setup

1. CUDA toolkit installation (done via laptop imaging)
 - a. Desktop
 - i. CUDA 10.2 with driver
 - ii. Disable VSync and Flipping in Nvidia OpenGL settings for best perf



2. Extract/git clone <insert tarfile name> lab files.
3. Install required packages

```
$ sudo apt-get install default-jre python libglew-dev freeglut3-dev
```
4. Install (download and unpack) Eclipse 4.9, e.g. from https://download.eclipse.org/technology/epp/downloads/release/2018-09/R/eclipse-cpp-2018-09-linux-gtk-x86_64.tar.gz
5. Install the Nsight Eclipse Edition eclipse plugin
 - a. Desktop: by running <cuda install dir>/bin/nsight_ee_plugins_manage.sh install <eclipse-path>
6. Generate the Nsight Eclipse Edition projects and launch configurations
 - a. Create clean directory /home/nvidia/cuda-workspace
 - b. We use the generate_nsightee_projects.py script

```
generate_nsightee_projects.py --src <code location> --dst /home/nvidia/cuda-workspace --steps 7
```
 - c. This will create all the project steps, a launch configuration for each step, as well as copy the image sources to the workspace directory

```
$ ls -l ~/cuda-workspace/
drwxrwxr-x 2 nvidia nvidia 4096 Jan 10 16:24 img
-rw-rw-r-- 1 nvidia nvidia 1003 Jan 16 17:55 launch_step_0.launch
-rw-rw-r-- 1 nvidia nvidia 1003 Jan 16 17:55 launch_step_1.launch

-rw-rw-r-- 1 nvidia nvidia 1003 Jan 16 17:55 launch_step_3.launch
-rw-rw-r-- 1 nvidia nvidia 1003 Jan 16 17:55 launch_step_4.launch
-rw-rw-r-- 1 nvidia nvidia 1003 Jan 16 17:55 launch_step_5.launch
-rw-rw-r-- 1 nvidia nvidia 1003 Jan 16 17:55 launch_step_6.launch
-rw-rw-r-- 1 nvidia nvidia 1003 Jan 16 17:55 launch_step_7.launch
drwxrwxr-x 2 nvidia nvidia 4096 Jan 16 18:15 step_0
drwxrwxr-x 2 nvidia nvidia 4096 Jan 16 18:17 step_1

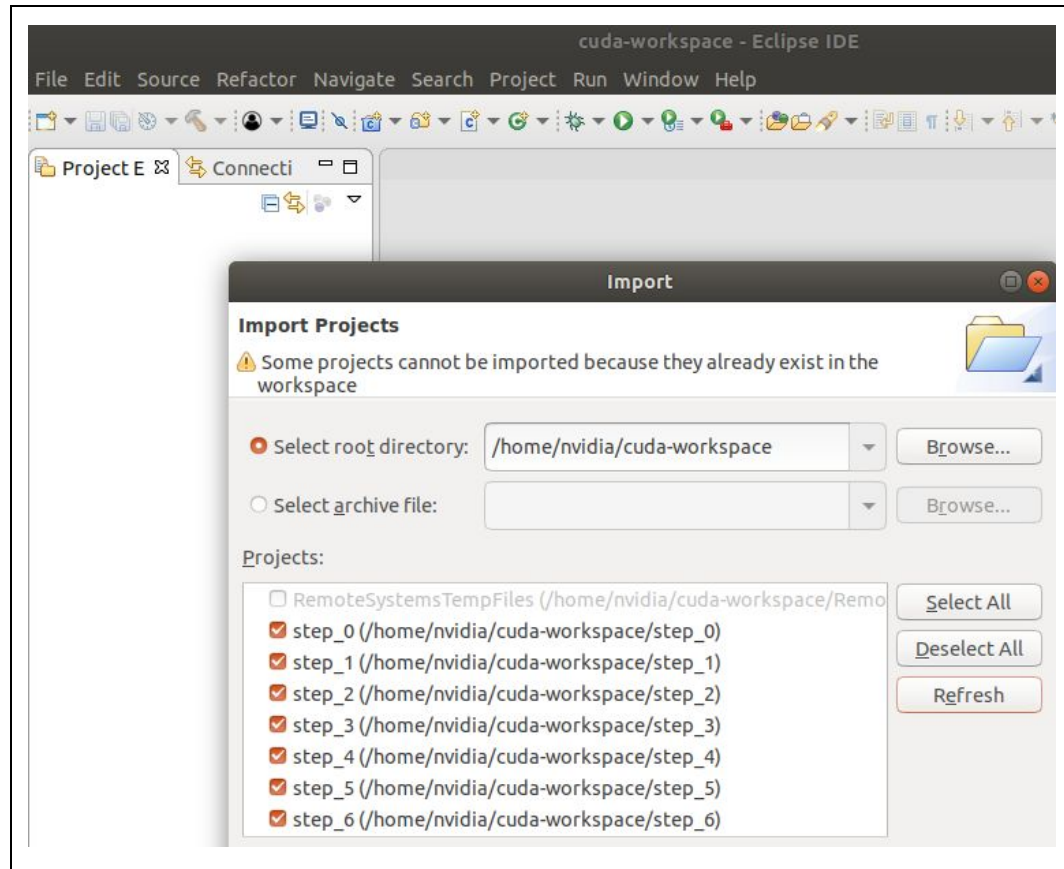
drwxrwxr-x 2 nvidia nvidia 4096 Jan 16 18:25 step_3
drwxrwxr-x 2 nvidia nvidia 4096 Jan 16 17:55 step_4
drwxrwxr-x 2 nvidia nvidia 4096 Jan 16 17:55 step_5
drwxrwxr-x 2 nvidia nvidia 4096 Jan 16 17:55 step_6
drwxrwxr-x 2 nvidia nvidia 4096 Jan 16 17:55 step_7
```

Step 0: Nsight Eclipse Edition (plugin)

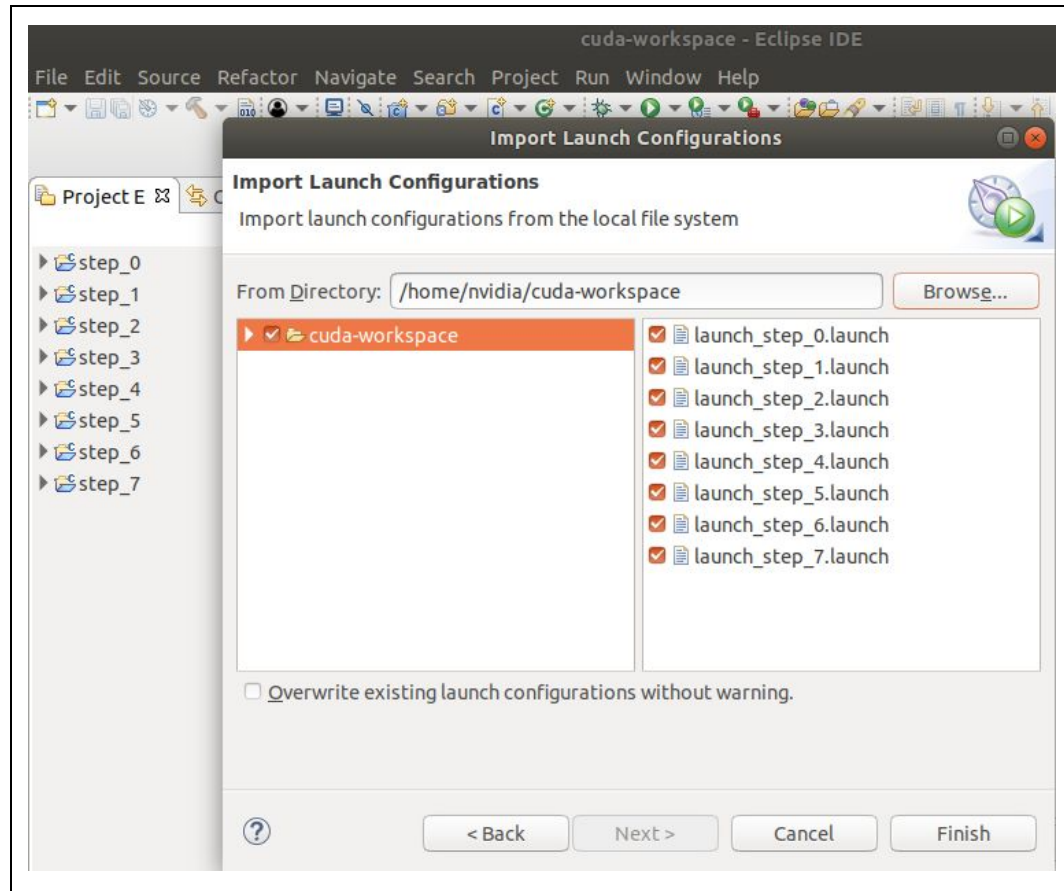
Step 0 Video: <https://developer.nvidia.com/gtc/2020/video/t21395-1>

Building the application using Nsight Eclipse Edition (plugin).

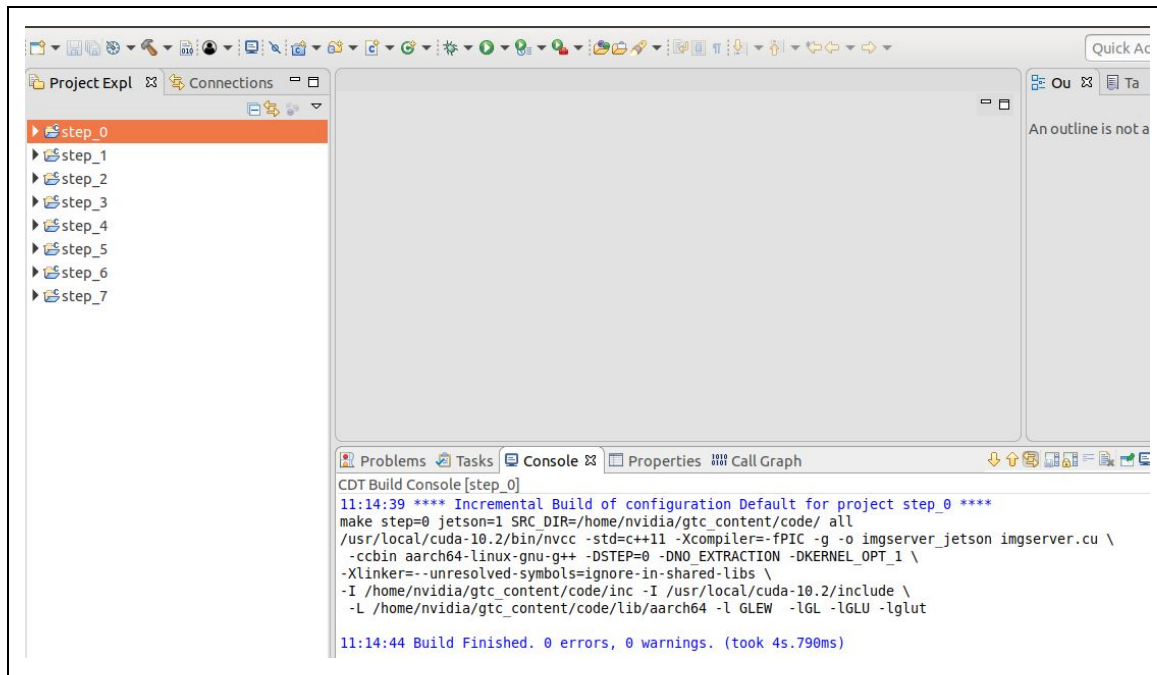
1. Open Eclipse, decide on Nsight Eclipse Edition plugin data collection preference
2. If asked, select ~/cuda-workspace as default workspace
3. Import projects from /home/nvidia/cuda-workspace
 - a. Project Explorer > Right Click > Import > General > Existing Projects into Workspace > Next
 - b. Specify /home/nvidia/cuda-workspace and select Refresh
 - c. Keep all projects enabled, note that step 2 will be missing
 - d. Finish



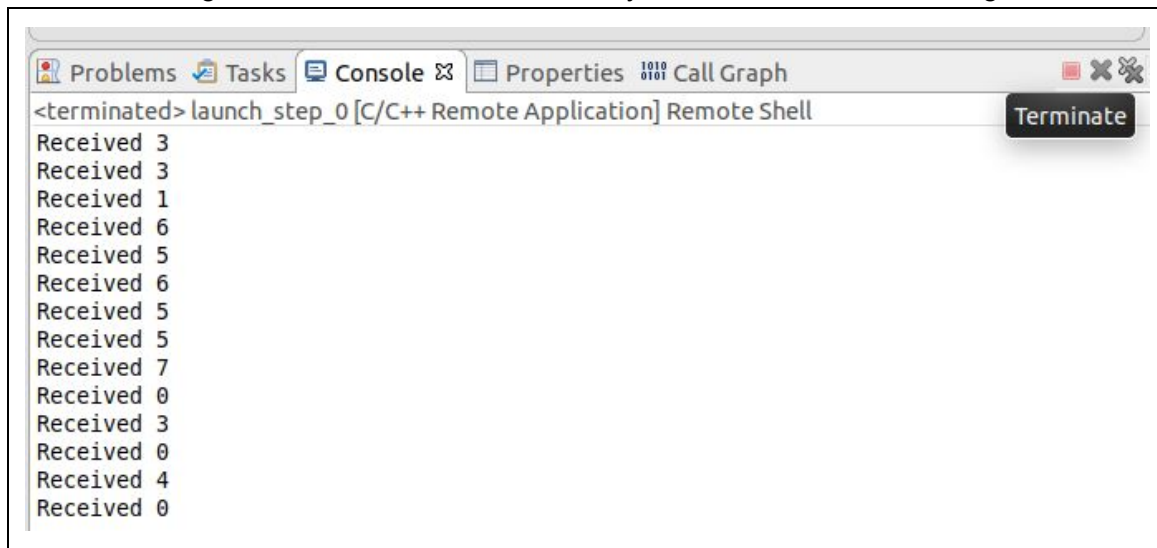
- e.
- 4. Import launch configs from /home/nvidia/cuda-workspace
 - a. Project Explorer > Right Click > Import > Run/Debug> Launch Configurations > Next
 - b. Browse /home/nvidia/cuda-workspace and select checkbox
 - c. Keep all launch configs enabled
 - d. Finish



5. Select step_0 and build (hammer icon)



6. Launch using the green Run button
7. Terminate using the red Terminate button once you see Received N messages



CLI Alternative:

1. Build the application with*

```
$ make clean  
$ make step=0
```

2. Run the app to see the anticipated app behavior (without any actual classification happening, all requests are simply rendered directly).

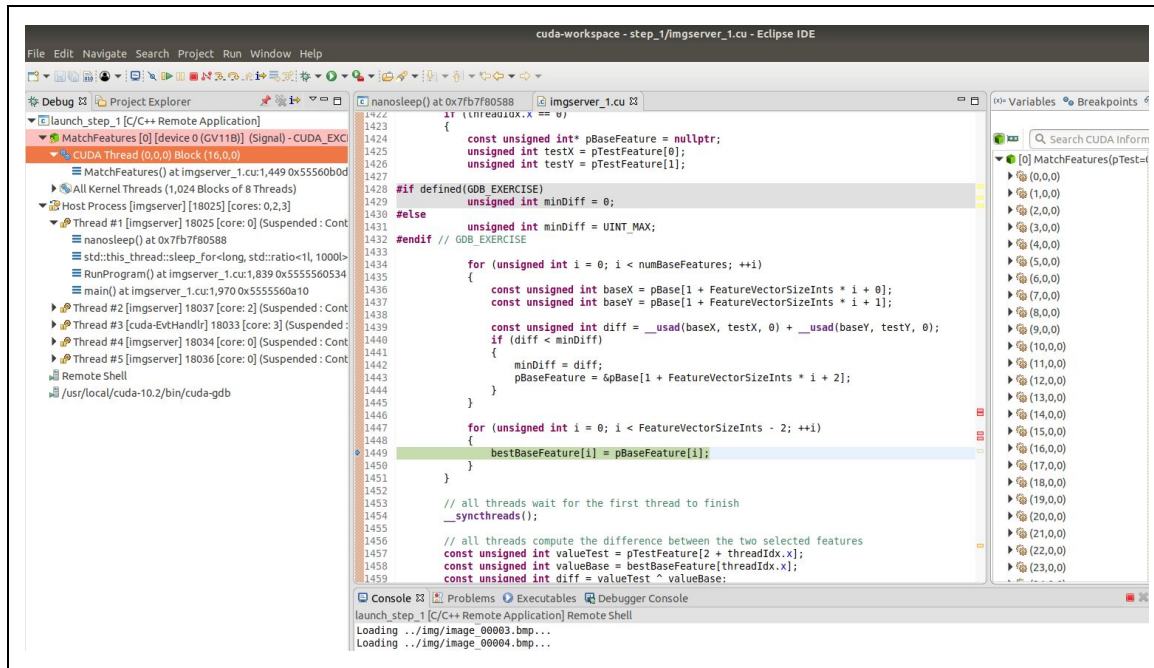
```
$ ./imgserver 50 100
```

Step 1: cuda-gdb

Step 1 Video: <https://developer.nvidia.com/gtc/2020/video/t21395-2>

Debugging a kernel crash using cuda-gdb.

1. Build the step_1 Nsight Eclipse Edition project using the hammer icon
2. Launch the project using the Run button (the connection had already been setup in step_0)
3. Watch the application crash and terminate in the Console window
4. Launch the project under cuda-gdb using the Debug button.
5. Wait until the application breaks in the kernel under the debugger. Then follow the same code inspection steps as in the CLI variant below.



6. Fix the initialization, recompile and Run again to see that the application now doesn't crash anymore.

CLI Alternative:

1. Build the application on the CLI with

```
$ make clean
$ make step=1
```

2. Run the application and observe kernel crash, caught by checkCudaErrors() macro, leading to application termination

```
$ ./imgserver 50 100
Exercise step 1
Loading ../FlowersDataset/bmp/image_00001.bmp...
Loading ../FlowersDataset/bmp/image_00002.bmp...
Loading ../FlowersDataset/bmp/image_00003.bmp...
Loading ../FlowersDataset/bmp/image_00004.bmp...
Loading ../FlowersDataset/bmp/image_00005.bmp...
Loading ../FlowersDataset/bmp/image_00006.bmp...
Loading ../FlowersDataset/bmp/image_00007.bmp...
Loading ../FlowersDataset/bmp/image_00008.bmp...
Loading ../FlowersDataset/bmp/image_00009.bmp...
Loading ../FlowersDataset/bmp/image_00010.bmp...
Loading smiley.bmp...
checkCudaErrors() API error = 0077 (an illegal memory access was
encountered) from file <imgserver.cu>, line 1518.
freeglut ERROR: Internal <Window Enumeration> function called without
first calling 'glutInit'.
checkCudaErrors() API error = 0029 (driver shutting down) from file
```

```
<imgserver.cu>, line 560.
```

3. Run the application under cuda-gdb until it breaks at the device-side kernel exception. For the CLI, see below. For Nsight Eclipse Edition, select Debug as Remote Application.

```
$ /usr/local/cuda-10.2/bin/cuda-gdb ./imgserver 50 100
NVIDIA (R) CUDA Debugger
10.2 release
...
Reading symbols from ./imgserver...done.
(cuda-gdb) r
Starting program:
/media/WORK/p4/sw/devtools/Presentations/2020/GTC/labs/code/imgserver
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
Exercise step 1
Loading ../FlowersDataset/bmp/image_00001.bmp...
...
Loading smiley.bmp...
[New Thread 0x7fffe7f0e700 (LWP 18699)]
...

CUDA Exception: Warp Illegal Address
The exception was triggered at PC 0x13fale0 (imgserver.cu:1411)

Thread 7 "imgserver" received signal CUDA_EXCEPTION_14, Warp Illegal Address.
[Switching focus to CUDA kernel 0, grid 12, block (141,0,0), thread (0,0,0),
device 0, sm 0, warp 1, lane 0]
0x00000000013falf0 in MatchFeatures<<<(1024,1,1),(8,1,1)>>>
(pTest=0x7fffbfe65400, pBase=0x7fffbfe00000, pResult=0x7fffbfe6f600) at
imgserver.cu:1411
1411                                     bestBaseFeature[i] = pBaseFeature[i];
(cuda-gdb)
```

4. Inspect the pBaseFeature variable, detect that it's set to 0x0. Inspect the code to see that it is set only under the (diff < minDiff) condition. Inspect minDiff to see that it is set (and initialized to) 0, so the condition will never evaluate to true. Inspect the initialization point to see that it must be fixed to minDiff = UINT_MAX.

```
(cuda-gdb) p pBaseFeature
$1 = (const unsigned int * @generic) 0x0
(cuda-gdb) p minDiff
$2 = 0
...
#if defined(GDB_EXERCISE)
    unsigned int minDiff = 0;
#else
    unsigned int minDiff = UINT_MAX;
#endif
    for (unsigned int i = 0; i < numBaseFeatures; ++i)
    {
        const unsigned int baseX = pBase[1 + FeatureVectorSizeInts
* i + 0];
        const unsigned int baseY = pBase[1 + FeatureVectorSizeInts
* i + 1];

        const unsigned int diff = __usad(baseX, testX, 0) +
```

```

__usad(baseY, testY, 0);
    if (diff < minDiff)
    {
        minDiff = diff;
        pBaseFeature = &pBase[1 + FeatureVectorSizeInts * i +
2];
    }

    for (unsigned int i = 0; i < FeatureVectorSizeInts - 2; ++i)
    {
        bestBaseFeature[i] = pBaseFeature[i];
    }

```

5. Fix the initialization, recompile using step=1 and run again to see that the application doesn't crash anymore.

Step 2: Compute Sanitizer*

* Compute Sanitizer is not available in CUDA 10.2. Skip this step if using CUDA 10.2.

Debugging unexpected application behavior using sanitizer's memcheck and initcheck tools.

1. Build the application with step=2 (or the respective Nsight Eclipse Edition project), this equals the fixed version of the debugger step.

```

$ make clean
$ make step=2

```

2. Run the application and note that almost all classification steps will fail.

```

$ ./imgserver
Exercise step 2
...
Loading ../FlowersDataset/bmp/image_00010.bmp...
Loading smiley.bmp...
Received 9, Classified as 9 - MATCH
Received 3, Classified as 7 - FAIL
Received 8, Classified as 1 - FAIL
Received 5, Classified as 3 - FAIL
Received 4, Classified as 3 - FAIL
Received 1, Classified as 9 - FAIL
Received 9, Classified as 7 - FAIL
Received 1, Classified as 9 - FAIL
Received 0, Classified as 1 - FAIL
Received 0, Classified as 3 - FAIL
Received 7, Classified as 7 - MATCH
Received 1, Classified as 6 - FAIL
Received 3, Classified as 0 - FAIL
Received 6, Classified as 1 - FAIL

```

```
Received 0, Classified as 2 - FAIL
Received 7, Classified as 3 - FAIL
...
```

3. Run the application again under compute-sanitizer, terminate the app after a few classification steps were reported. Note that sanitizer reports no errors.

```
$ compute-sanitizer ./imgserver
===== COMPUTE-SANITIZER
Exercise step 2
Loading ../FlowersDataset/bmp/image_00001.bmp...
...
Loading smiley.bmp...
Received 9, Classified as 8 - FAIL
Received 3, Classified as 8 - FAIL
Received 8, Classified as 1 - FAIL
Received 5, Classified as 4 - FAIL
Received 4, Classified as 5 - FAIL
Received 1, Classified as 4 - FAIL
Received 9, Classified as 1 - FAIL
<manually terminate application>
===== ERROR SUMMARY: 0 errors
```

4. Note that Compute Sanitizer has four sub-tools (memcheck (default), initcheck, syncheck and racecheck). Once memcheck doesn't report any errors, we can use the other three tools. We don't know which one to start with, but for the sake of time, we can start with initcheck, which will show the prepared error.
5. Run the application again using Compute Ssanitizer's initcheck tool until errors are reported, then terminate the app (e.g. by pressing Ctrl-C) and inspect the errors.

```
$ compute-sanitizer --tool initcheck ./imgserver
===== Uninitialized __global__ memory read of size 4 bytes
===== at 0x140 in
/media/WORK/p4/sw/devtools/Presentations/2020/GTC/labs/code/imgserver.cu:1388:
MatchFeatures(unsigned int const *,unsigned int const *,unsigned int*)
===== by thread (0,0,0) in block (367,0,0)
===== Address 0x7f4a17a68d5c
===== Saved host backtrace up to driver entry point at kernel launch
time
===== Host Frame:cuLaunchKernel [0x7f4a4713b546]
===== in /usr/lib/x86_64-linux-gnu/libcuda.so
===== Host Frame:[0x443312]
===== in
/media/WORK/p4/sw/devtools/Presentations/2020/GTC/labs/code/./imgserver
===== Host Frame: [0x443507]
===== in
/media/WORK/p4/sw/devtools/Presentations/2020/GTC/labs/code/./imgserver
...
===== Host Frame: [0x7f4a4a313c80]
===== in /usr/lib/x86_64-linux-gnu/libstdc++.so.6
===== Host Frame: [0x7f4a4a7e86ba]
===== in /lib/x86_64-linux-gnu/libpthread.so.0
===== Host Frame:clone [0x7f4a49d8241d]
===== in /lib/x86_64-linux-gnu/libc.so.6
=====
===== Internal Sanitizer Warning: Error buffer overflow has been detected.
Some records have been dropped
=====
```

```
Received 9, Classified as 9 - MATCH
===== ERROR SUMMARY: 2570 errors
```

6. Inspect the referenced code location (exact line may change from what is listed in this document). Apparently the feature we read from is not properly initialized. From the comment, we see which feature entry is selected by each block.

```
// from the vector, select the feature assigned to this block
//          vvvvv
// [<num>|..0..|..1..|..2..|..3..|..numTestFeatures-1..]
// <--->          <--->
// sizeof(unsigned int) sizeof(int) * FeatureVectorSizeInts
const unsigned int* pTestFeature = &pTest[1 + FeatureVectorSizeInts *
blockIdx.x];
__shared__ unsigned int bestBaseFeature[FeatureVectorSizeInts - 2];

// the first thread in the block searches the best-matching feature in
the base features vector
if (threadIdx.x == 0)
{
    const unsigned int* pBaseFeature = nullptr;
    unsigned int testX = pTestFeature[0];
```

7. We need to inspect the location where the features are written, which is in the ExtractFeatures kernel. In the code and comment, we see that the written entry is supposed to be selected based on the fvi variable (feature vector index), but we failed to multiply the entry size with this index, so we always write to only the second feature, for all blocks.

```
// atomically select the next feature vector index (fvi)
const unsigned int fvi = atomicAdd(&pFeatureVectors[0], 1);
if (fvi >= MaxFeatureVectors)
{
    // no more features available, mark this in the output image
    pOut[outIdx] = make_uchar4(0, 255, 0, 255);
}
else
{
    // compute the feature vector for this keypoint

    // we want to use the same pseudo-random number sequence for
each keypoint,
    // so we initialize the cuRand state every time with the same
seeds

    curandState_t randState;
    curand_init(0, 0, 0, &randState);

    // keypoints are randomly selected in a 65x65 search are
constexpr int patchRadius = 32;

    // from the vector, select the feature assigned to this block
    //          fvi
    //          vvvvv
    // [<num>|..0..|..1..|..2..|..3..|..numTestFeatures-1..]
    // <--->          <--->
    // sizeof(unsigned int) sizeof(int) * FeatureVectorSizeInts
#if defined(GDB_EXERCISE) || defined(SANITIZER_EXERCISE)
    unsigned int* pFeature = &pFeatureVectors[1 +
```

```

FeatureVectorSizeInts];
#else
        unsigned int* pFeature = &pFeatureVectors[1 +
FeatureVectorSizeInts * fvi];
#endif

```

8. After fixing the bug and re-compiling using step=2, we see that initcheck reports no errors anymore and classification is working as expected.

```

compute-sanitizer --tool initcheck ./imgserver
===== COMPUTE-SANITIZER
Exercise step 3
Loading ../FlowersDataset/bmp/image_00001.bmp...
...
Loading smiley.bmp...
Received 9, Classified as 9 - MATCH
Received 3, Classified as 3 - MATCH
Received 8, Classified as 8 - MATCH
Received 1, Classified as 1 - MATCH
Received 5, Classified as 5 - MATCH
Received 1, Classified as 1 - MATCH
Received 9, Classified as 9 - MATCH
Received 4, Classified as 4 - MATCH
===== ERROR SUMMARY: 0 errors

```

9. We can now run the app without Compute Sanitizer to see the same working behavior. However, it's still too slow, so many user requests are over the time threshold, resulting in sad smileys rather than flower images. This leads us to the performance analysis part of the training.

Steps 3-5: Nsight Systems

Step 3-4 Video: <https://developer.nvidia.com/gtc/2020/video/t21395-3>

Step 3: Analyze original application behavior

1. Build the first application step without functional bugs

```

$ make clean
$ make step=3

```


2. Set the correct kernel perf event paranoid flag, to collect as much CPU data as possible

```

$ sudo sh -c 'echo 1 > /proc/sys/kernel/perf_event_paranoid'

```

3. Start Nsight Systems and open a new project. Connect to your localhost.



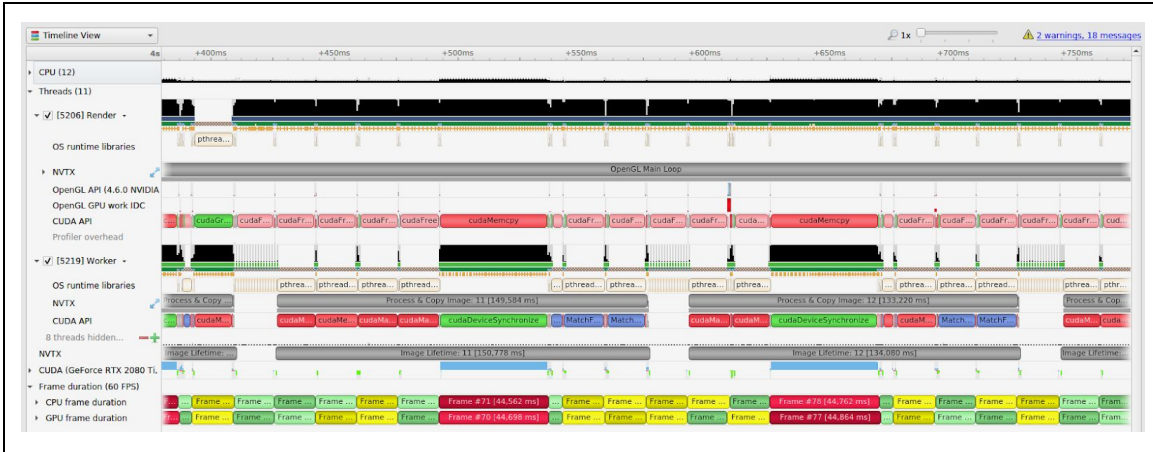
4. Once connected, fill in the project details.
 - a. Enable Collect Call stacks of executing thread, CUDA, OpenGL, NVTX and OS lib trace enabled,
 - b. Fill in target application and working dir
 App: /home/nvidia/cuda-workspace/step_3/imgserver 70 100
 WD: /home/nvidia/cuda-workspace/step_3

5. Launch data collection with Start and inspect the resulting trace. Once the application ends, stop data collection.

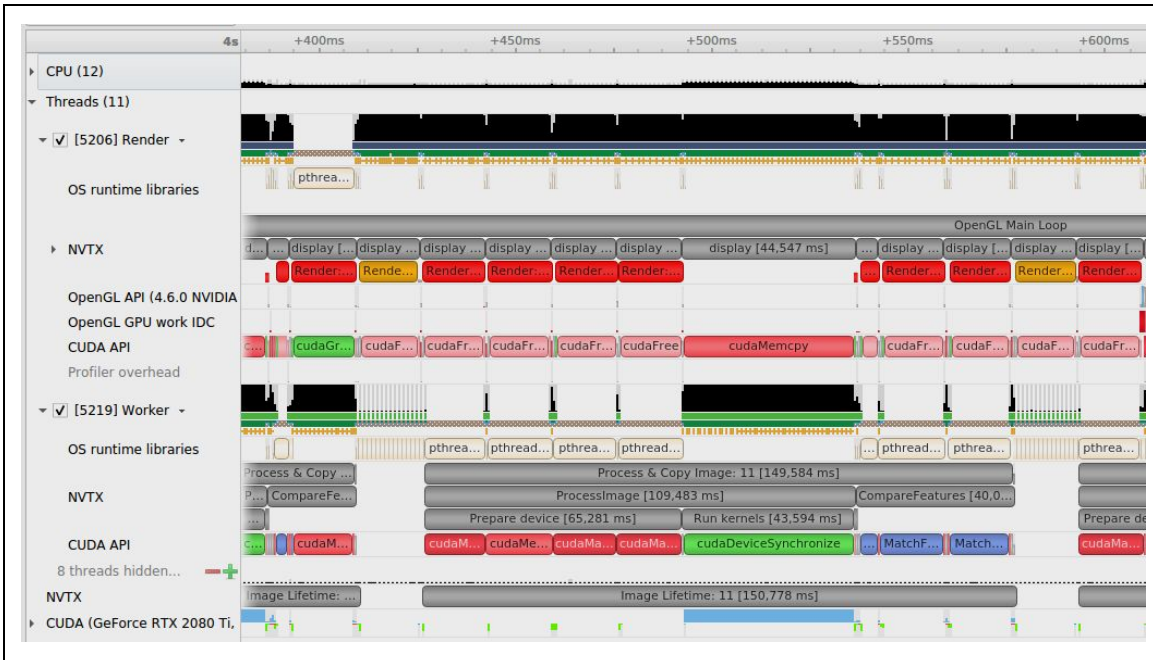
During data collection, most requests should exceed the threshold, resulting in sad smiley renderings. The trace will show 2-3 threads (Worker and Render) in three phases: Loading images from disk, Compute Features DB and the actual request processing. We are interested in optimizing the last phase (OpenGL Main Loop).



6. Zoom into the last phase until you see the Process & Copy image patterns. We see kernel execution (light blue) on the GPU, but we also see that the Render thread spends a considerable amount of time in cudaFree and cudaMemcpy API calls.



7. Zooming in further into the Render thread and expanding the NVTX rows provides more details: While the next request is processed on the Worker thread, the Render thread continuously calls `display()` which tries to Render to the output grid. During each Render call, we seem to free (and allocate, when zooming in further) device memory.



8. It would be smarter to allocate this buffer only once, and free it once all requests are done, since all request images are of the same size in our application. This is done for this buffer (and several others) in the next step.

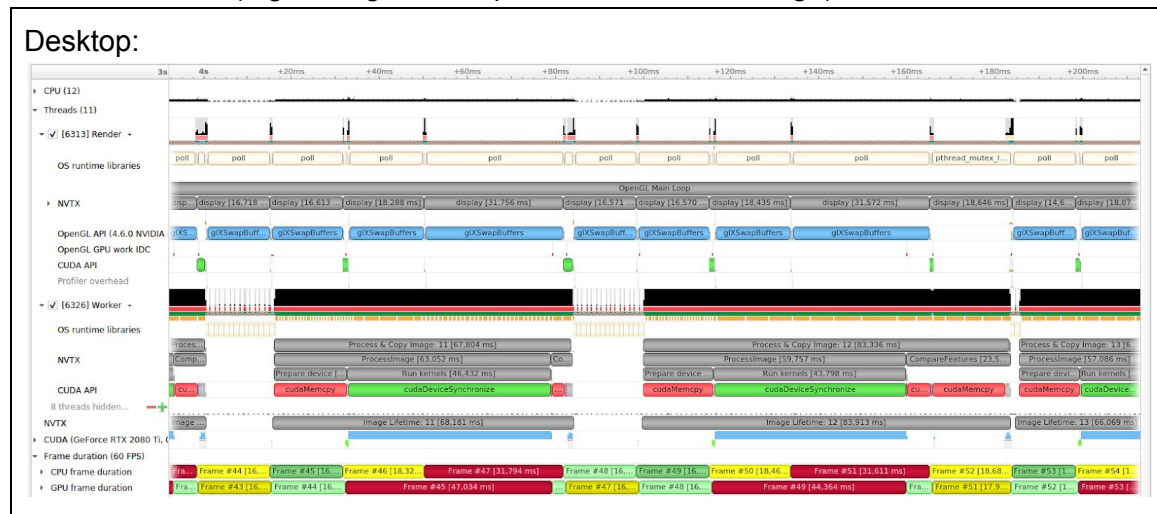
Step 4: Buffer reuse to minimize cudaMalloc/cudaFree

1. Build the next step, which enables device buffer reuse

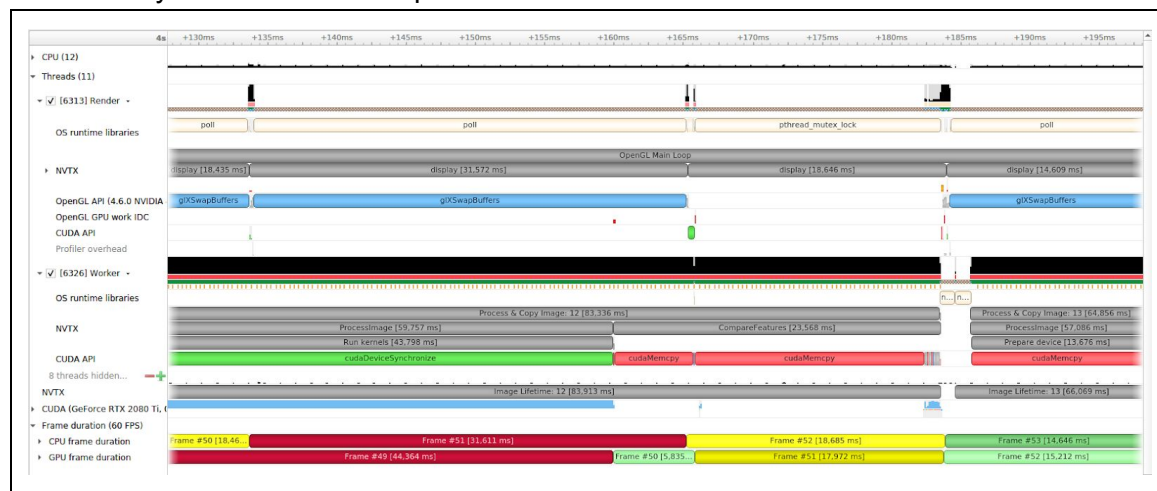
```
$ make clean
$ make step=4
```

2. Repeat the steps from the previous step 3 to collect a new trace in Nsight Systems. If Nsight Systems is still open, simply select the same Project and click Start. Then inspect

the collected trace. We see that the cudaFree calls are gone and so are several cudaMalloc calls (e.g. during the “Prepare device” NVTX range).



- Even though the Image Lifetime NVTX ranges became faster, we still see the large gaps between the processing phases, and we still see the Render thread calling display() continuously. There are various CUDA and OpenGL API calls happening, many of which can interfere with our Worker thread. For example, in the screenshot we see that several cudaMemcpy calls are extremely long even though they transfer the same amount of data as very small ones in that phase.



- The solution to this is simple: since our application is a very strict pipeline (Source request, Worker processing, Render output), there is no need for the Render thread to become active until the Worker thread is done processing the request and adding it to the transfer queue.

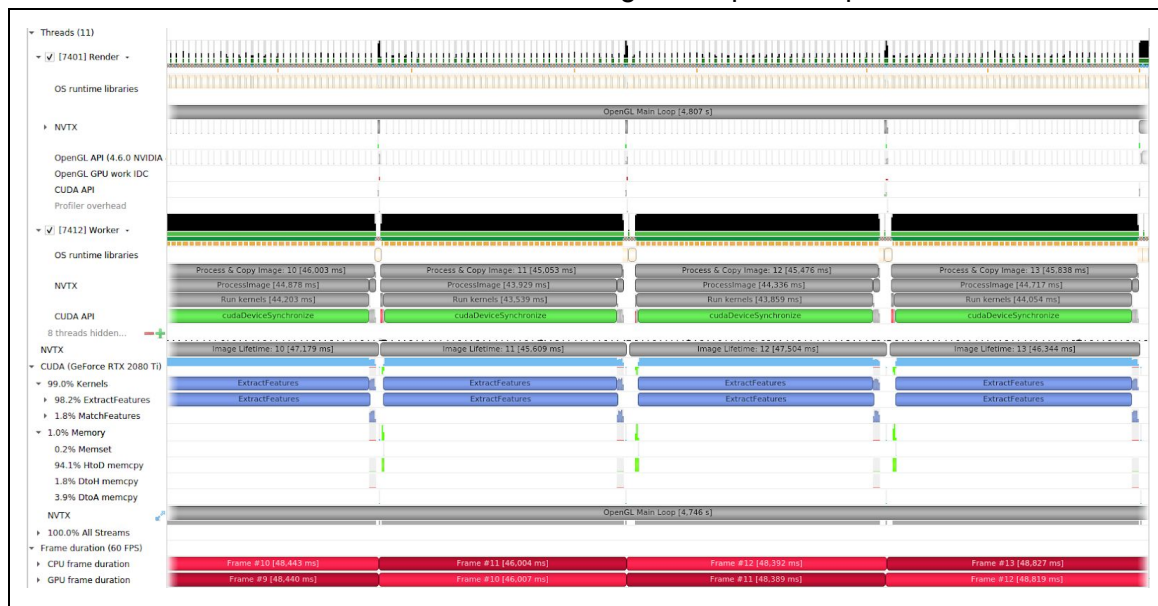
Step 5: Rendering on Demand

Step 5-7 Video: <https://developer.nvidia.com/gtc/2020/video/t21395-4>

1. Build the next step, which updates the Render thread to become active only when a new request was processed

```
$ make clean  
$ make step=5
```

5. Repeat the steps from the previous step 4 to collect a new trace in Nsight Systems. If Nsight Systems is still open, simply select the same Project and click Start. During the capture, we should now see some of the requests being served in time, which has them render the proper image to the output window. Afterwards, inspect the collected trace.
6. We see that each iteration is now executing much more regularly. Also, the Render thread only becomes active between two Process & Copy ranges. The average frame duration reported by Nsight Systems (which is somewhat equal to our Image Lifetime NVTX range) is ~50ms. The ExtractFeatures kernel is now the clear performance bottleneck. In the next exercise, we will use Nsight Compute to optimize this kernel.



Steps 6-7: Nsight Compute

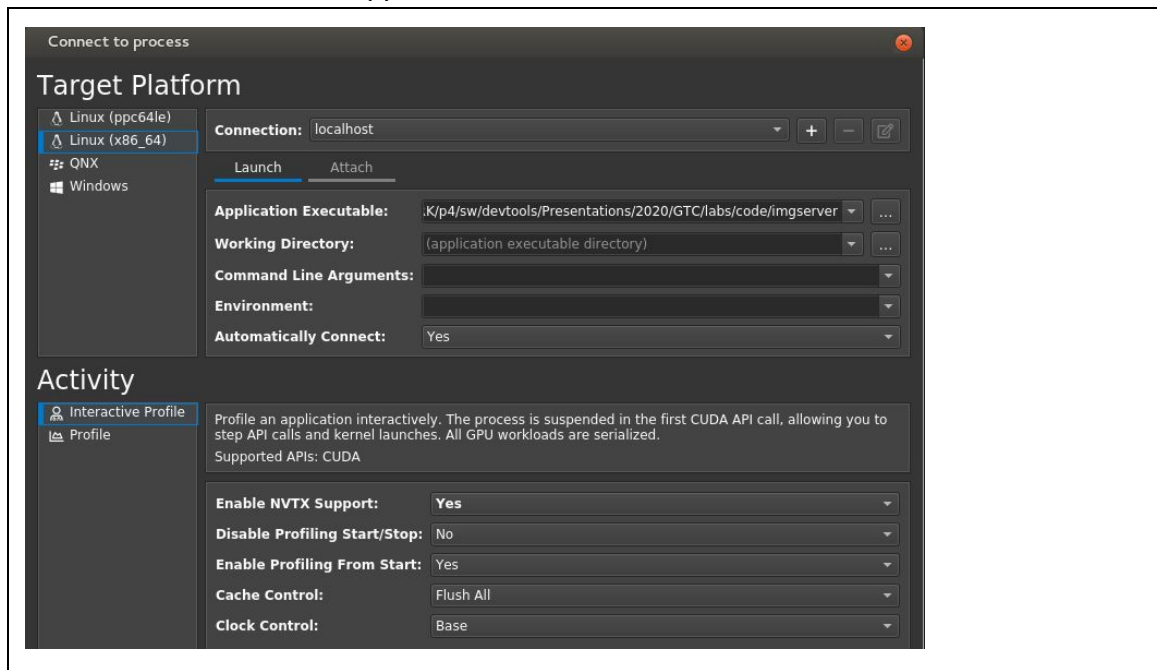
Analyzing and optimizing kernel performance using Nsight Compute interactive profiling.

Step 6.1: Analyze ExtractFeatures kernel

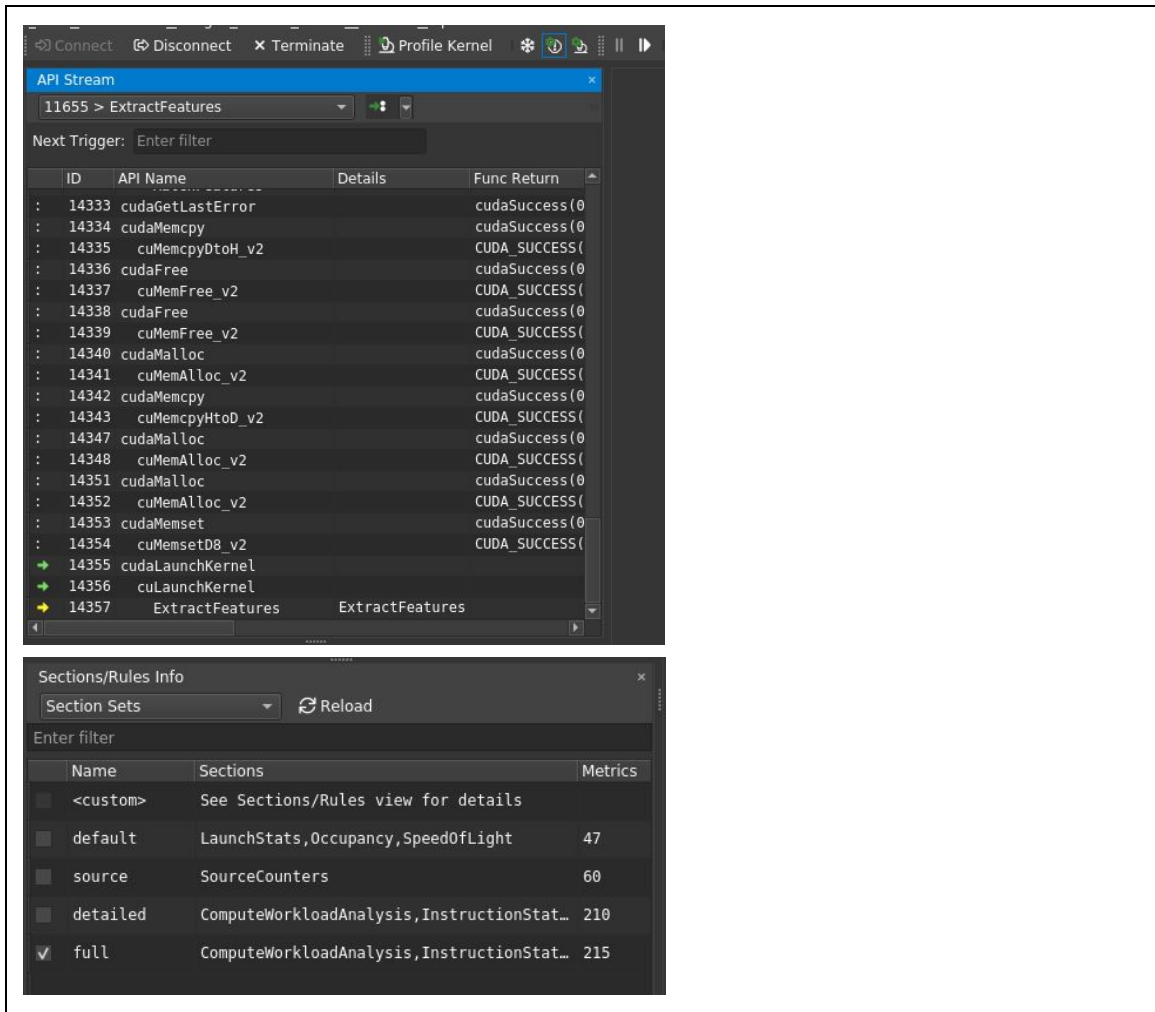
1. Re-use the application built as part of the last Nsight Systems exercise, or rebuild it using:

```
$ make clean  
$ make step=5
```


2. Nsight Systems has shown that the ExtractFeatures kernel is the primary bottleneck, so we start by analyzing that kernel. Start the Nsight Compute UI. In CUDA 10.2, the tool is called `nv-nsight-cu`. In CUDA 11 it is `ncu-ui`. Open a new project file and fill in the Application Executable to point to `step_5/imgserver`. Set Command Line Arguments to “50 100”, then Launch the application.



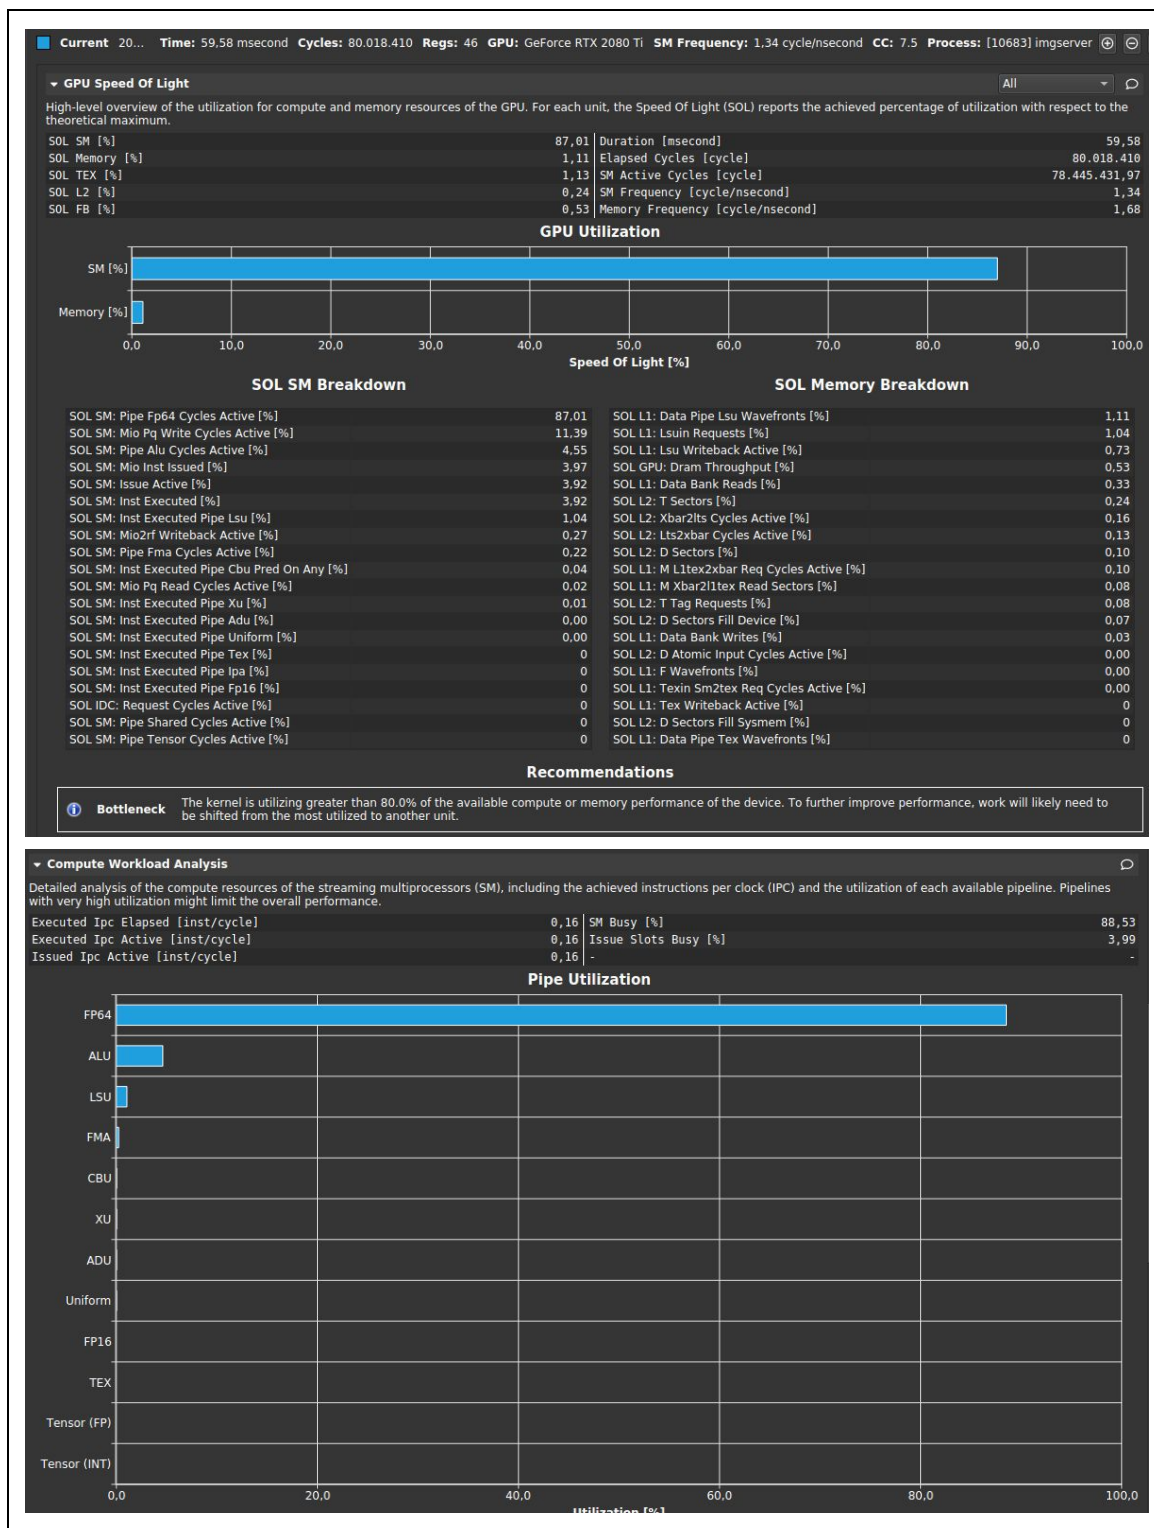
3. Once connected, the app is suspended in the first CUDA API call. Select “Run to Next Kernel” until you are at an ExtractFeatures kernel (it should be the very first one). Ensure that the “full” section set is enabled before selecting “Profile Kernel”.



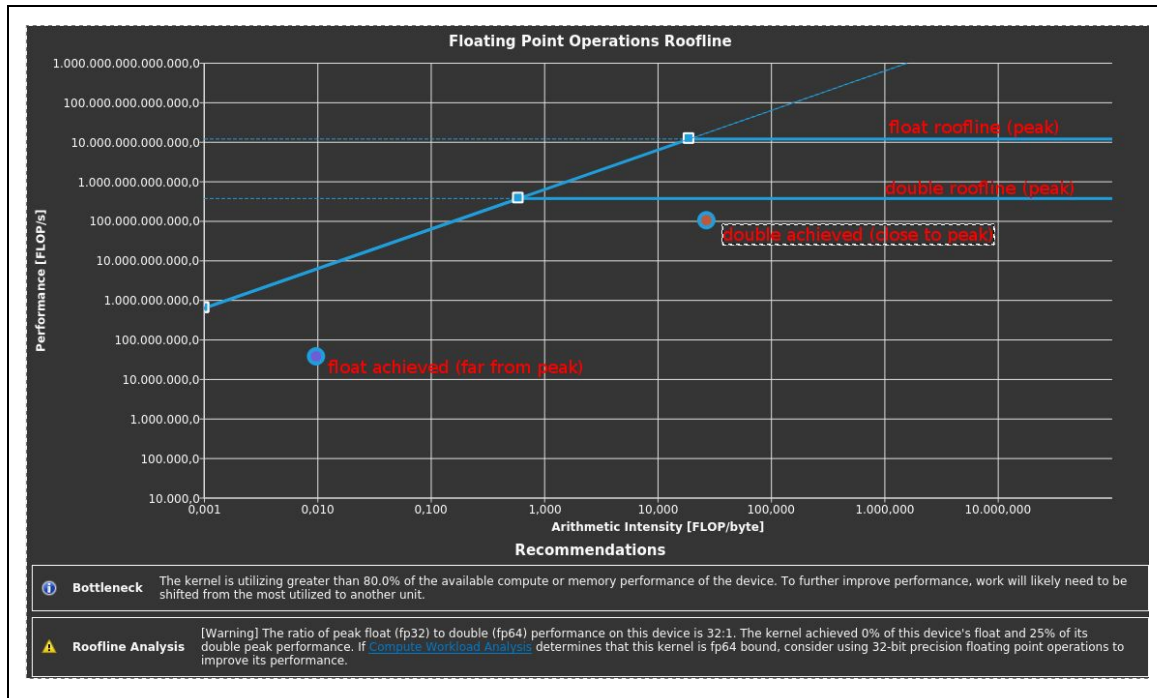
- After profiling is done and the report was created, we can Terminate the target application and start analyzing the collected report.



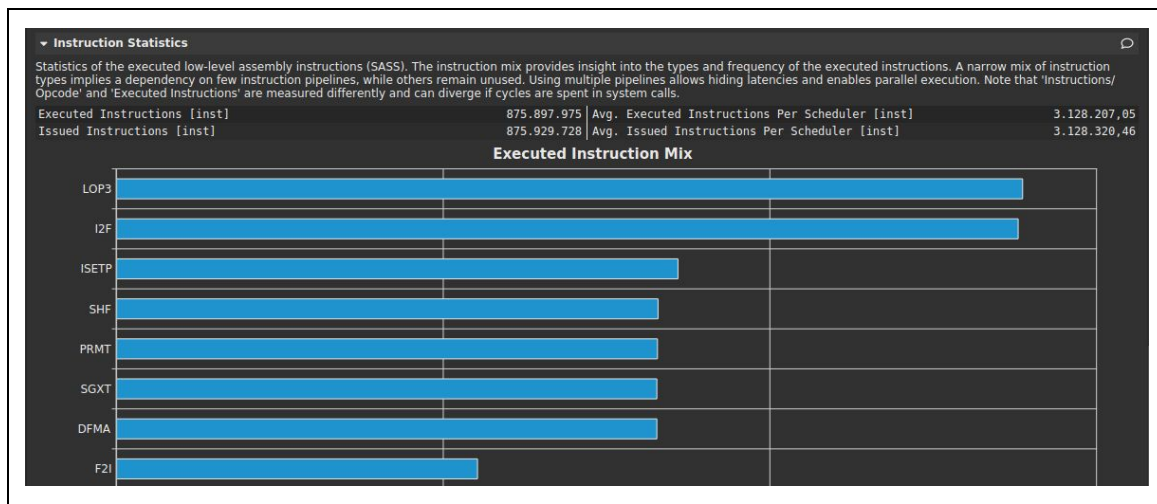
- We can see that the kernel is heavily compute bound, utilizing the SM units > 80%. We can see from the SOL SM Breakdown that the Fp64 (64bit/double floating point math) pipeline is by far the biggest contributor. Similar data can be found in the Compute Workload Analysis section



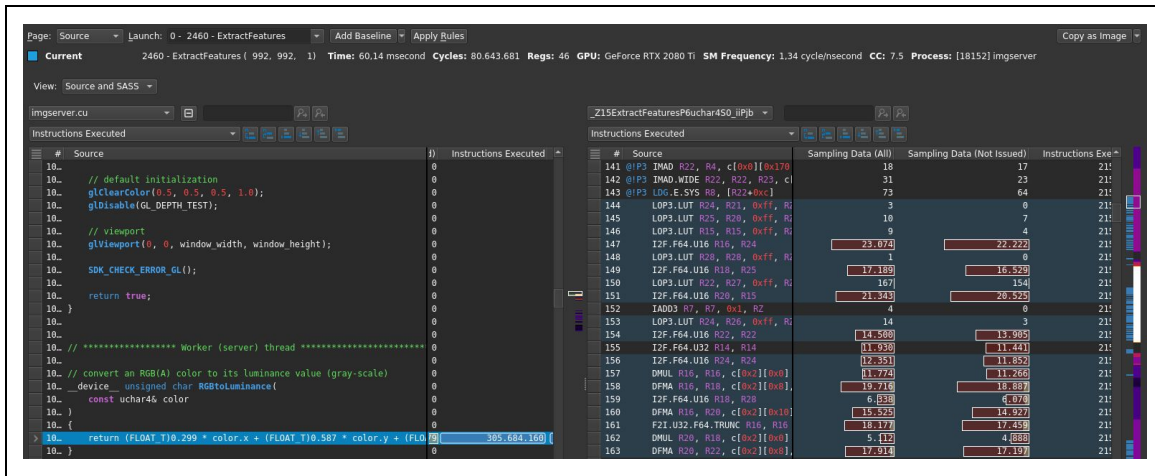
- If using Nsight Compute from CUDA 11.0, the roofline chart guides the user to use float instead of double due to the much 32:1 peak perf ratio between 32 and 64bit units. Red labels added in screenshot for explanatory purposes.



7. We can use the Instruction Statistics section to learn more about the types and counts of instructions executed, which show a large number of floating point related instructions for this kernel



8. Switching to the Source page using the Page dropdown on the top of the report, we can see where in the code those instructions are executed. For that, make sure that "Instructions Executed" is selected in the metrics drop down and then use the "Move to the row with the highest value" button from the navigation buttons. We see that most instructions are executed as part of the RGB to luminance/grayscale conversion



- On the left side, we see the CUDA C code, with the correlated SASS highlighted on the right side. Checking the definition of `FLOAT_T` in our code, we see that it is currently defined to double, which is why we do 64bit/double math right now.

```
#if defined(KERNEL_OPT_1)
    #define FLOAT_T double
#else
    #define FLOAT_T float
#endif
```

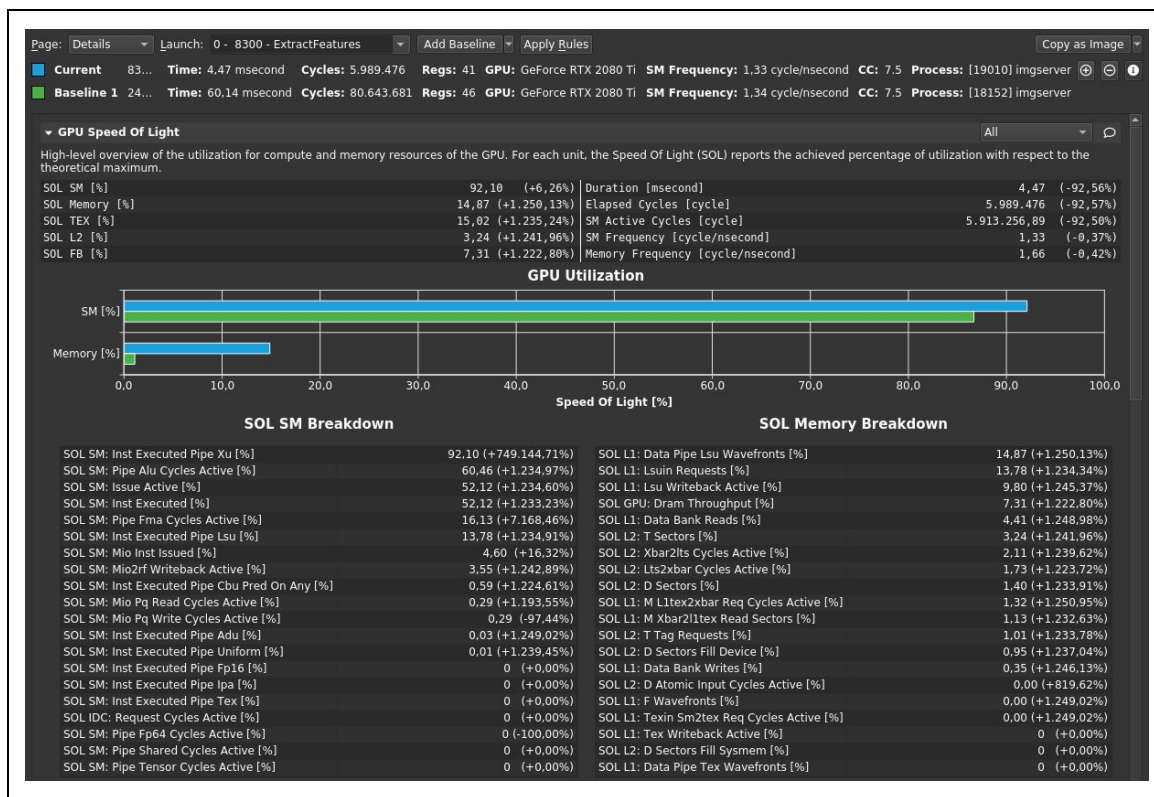
- Since we don't really need double-precision, we can use 32bit/float math instead, given that we have a 32:1 ratio for float to double units on Turing RTX GPUs (for desktop, the user is notified of this by the roofline analysis).

Step 6.2: Analyze the 32bit float math ExtractFeatures kernel

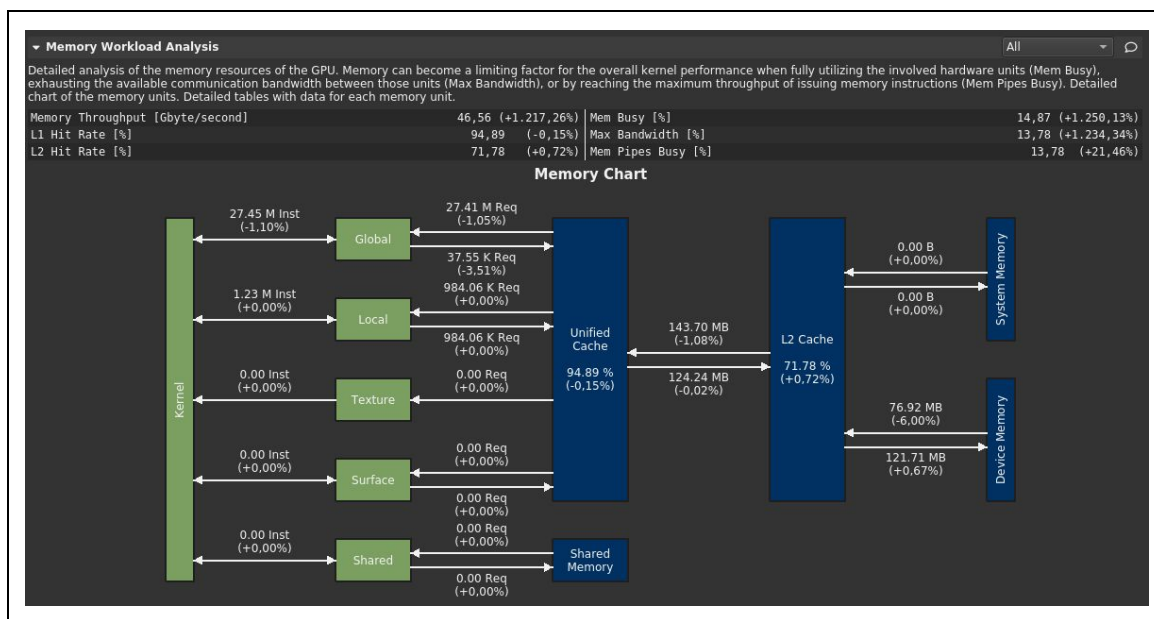
- Build the optimized kernel using 32bit instead of 64bit math using:

```
$ make clean
$ make step=6
```

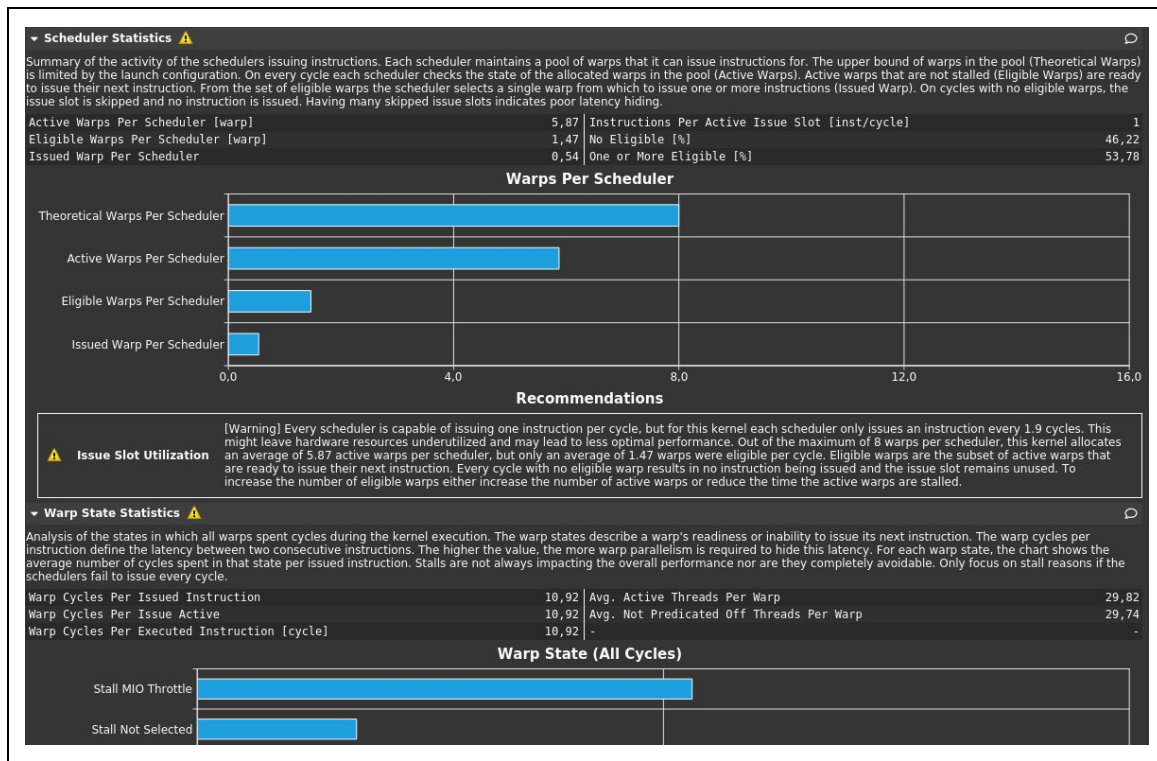
- Then, repeat the steps from the previous exercise 6.1 to collect the data for the optimized ExtractFeatures kernel in a second report.
- We can briefly go back to the first report, select "Add Baseline", and then switch to the current once again. We now see all data from the first, unoptimized kernel compared against the Current (blue) kernel. Most notably, it is 92% faster (Duration), and doesn't use 64bit (Fp64) math anymore. Since we are blocked less by the fp64 pipeline, we now have a 12x better memory units utilization.



- Inspecting the Memory Workload Analysis section, we see that we have a high cache hit rate in both L1 and L2, indicating that we re-read the same global memory multiple times.



- Analyzing the two sections with analysis recommendations, Scheduler and Warp State Statistics, we see that we issue only every second cycle, instead of every cycle, and that we are primarily stalled by MIO Throttle, which can also hint towards math instructions (see rule message).



- Moving to the Source page and selecting "Sampling Data (All)" from the metrics drop down, we can see when navigating to the top N stall locations that most of them are within the functions reading pixels, converting them to luminance and applying gaussian blur.

#	Source	Sampling Data (All)	Sampling Data (Not Issued)	Instructions Executed	Predicated-On Thread Instructions Executed	Memory Address Space	Memory*
10..	}	0	0				
10..	{	0	0				
10..	return (FLOAT_T)0.299 * color.x + (FLOAT_T)0.587 * color.y + (FLOAT_T)0.099 * color.z;	87,984	41,440	362,467,840	9,668,324,764		
10..	}	0	0				
10..	_device_ uchar4 GetPixel(10.. uchar4* pimg, 10.. int x, 10.. int y, 10.. int imgw, 10.. int imgh)	0	0				
11..	{	0	0				
11..	if (x >= 0 && y >= 0 && x < imgw && y < imgh)	2,448	7,544	64,329,504	4,924,409,280		
11..	{	0	0				
11..	return pimg[y * imgw + x];	24,999	9,612	193,683,840	5,777,328,168	Global(29)	
11..	}	0	0				
11..	return {0, 0, 0, 0};	0	0				
11..	}	0	0				
11..	#if defined(KERNEL_OPT_1) defined(KERNEL_OPT_2)	0	0				
11..	{	0	0				
11..	_device_ unsigned char GetLuminance(11.. uchar4* pimg, 11.. int x, 11.. int y, 11.. int imgw, 11.. int imgh, 11.. bool applyBlur)	0	0				
11..	{	0	0				
11..	if (!applyBlur)	14	0	621,472	1,968,128		
11..	{	0	0				
11..	return RGBtoLuminance(GetPixel(pimg, x, y, imgw, imgh));	0	0				
11..	}	0	0				
11..	else	0	0				
11..	{	0	0				
11..	unsigned int value = 0;	0	0				
11..	const int radius = 3;	0	0				
11..	const auto num = (radius + 1) * (radius + 1);	0	0				
11..	for (int dy = -radius; dy <= radius; dy++)	1,301	408	20,158,848	587,638,800		
11..	{	0	0				
11..	for (int dx = -radius; dx <= radius; dx++)	0	0				
11..	{	0	0				
11..	value += RGBtoLuminance(GetPixel(pimg, x + dx, y + dy, imgw, imgh));	57,987	28,023	171,305,184	5,143,844,604		
11..	}	0	0				
11..	}	0	0				

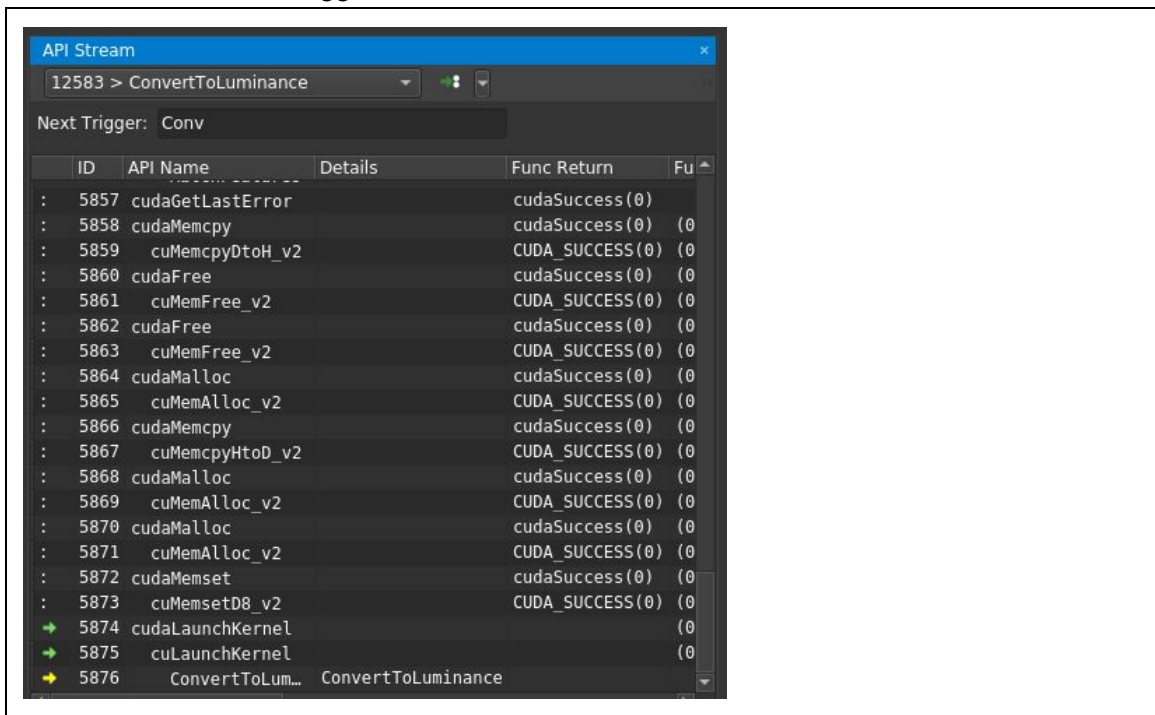
- The current implementation of the kernel uses an RGB input image and does all luminance and blur computations “live”, i.e. whenever a pixel is read. If pixels are read multiple times, because they are part of multiple keypoint feature extractions for example, all computations will be done multiple times. This can be made more efficient by extracting the grayscale conversion and gaussian blur steps into separate kernels, run prior to the actual feature extraction kernel. Then, both grayscale conversion and blur computation are done exactly once per pixel.

Step 7: Analyze the three optimized kernels

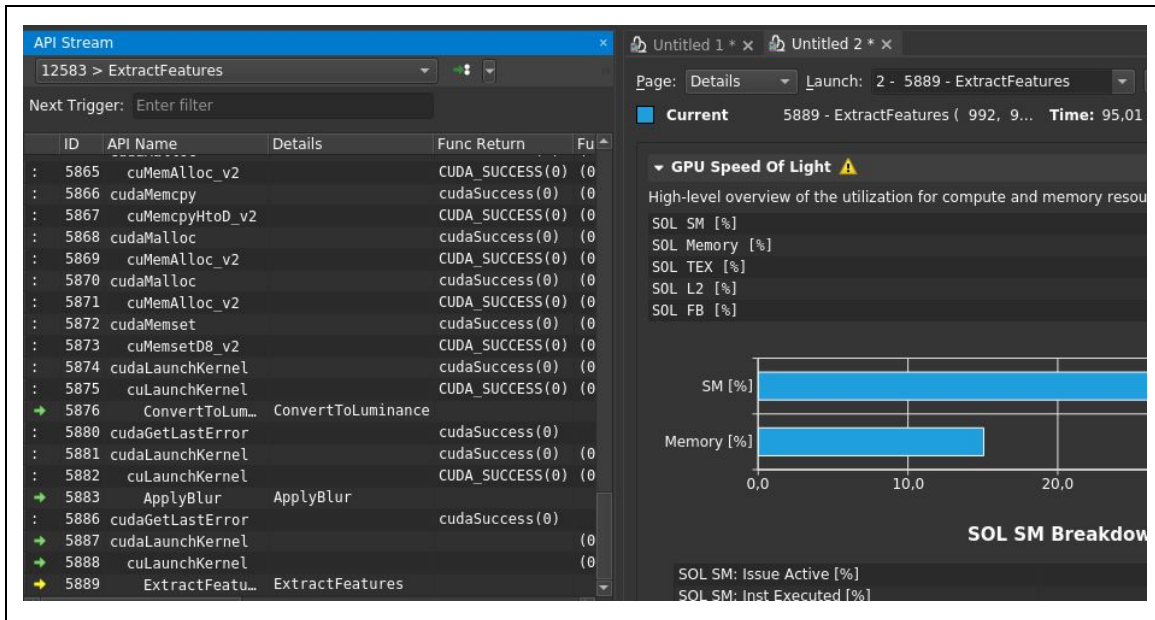
- Build the optimized application using three consecutive kernels per image with

```
$ make clean
$ make step=7
```

- Connect to the application again using the Interactive Profile activity. Once more, Resume the application to a steady state, and Pause.
- We want to analyze three matching ConvertToLuminance, ApplyBlur and ExtractFeatures kernels. If not already suspended at a ConvertToLuminance kernel, insert “Conv” into the Next Trigger edit and select “Run to Next Kernel”. Then remove “Conv” from the Next Trigger edit.



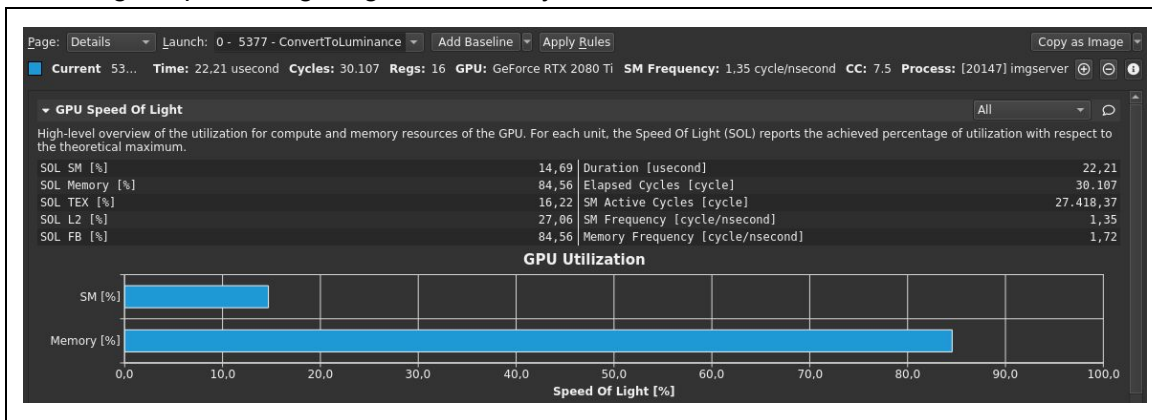
- Use the “full” section set to profile this kernel, then two times advance to the next kernel using “Run to Next Kernel” and profile both the ApplyBlur and the ExtractFeatures kernels using the same set.

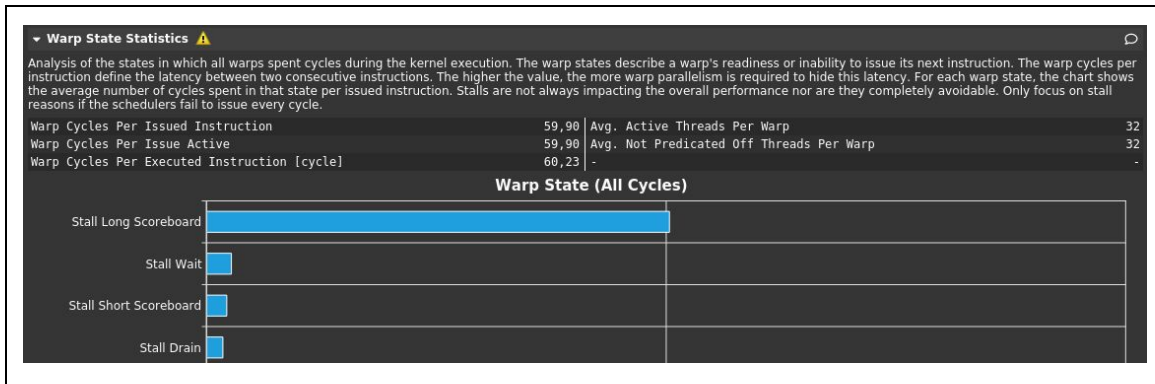


- When done, Terminate the target application and inspect the collected report.
- Switching to the Summary page, we can quickly see that the runtime (in cycles) of the three kernels combined is much smaller than the previous runtime of ExtractFeatures (~200k cycles vs ~6000k cycles). We cannot do a useful direct kernel-to-kernel comparison anymore, since the workload from ExtractFeatures got split up across the three new kernels.

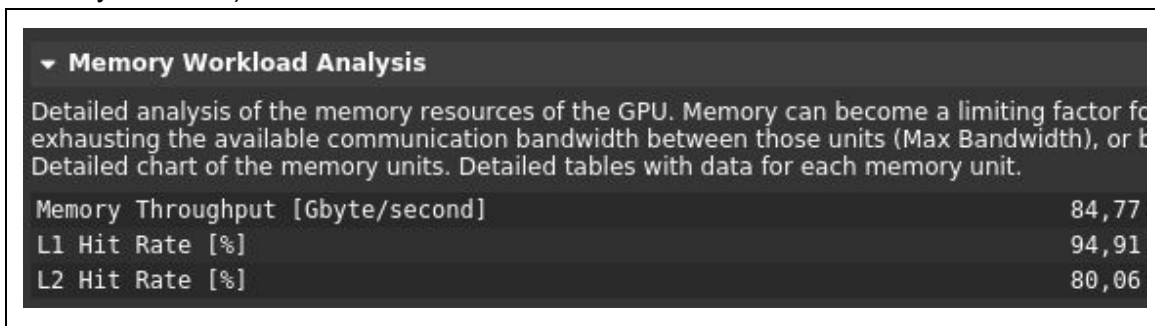
Function Name	Demangled Name	Process	Device Name	Cycles [cycle]	SOL Memory [%]	SOL SM [%]
ConvertToLumin...	ConvertToLum...	[20147] img...	GeForce RTX ...	30.107	84,56	14,69
ApplyBlur	ApplyBlur(uc...	[20147] img...	GeForce RTX ...	139.290	37,26	93,56
ExtractFeatures	ExtractFeatu...	[20147] img...	GeForce RTX ...	114.888	16,57	55,97

- Back on the Details page, we can use the Launch drop down to switch to the first profiled kernel, ConvertToLuminance. We see that it is memory-bound, utilizing the memory units to ~85%. The same conclusion can be drawn from the Compute Workload Analysis section, which shows that LSU (Load Store Unit) is the most-utilized pipeline, and in the Warp State Statistics, which is clearly dominated by Long Scoreboard Stalls, indicating warps waiting on global memory.

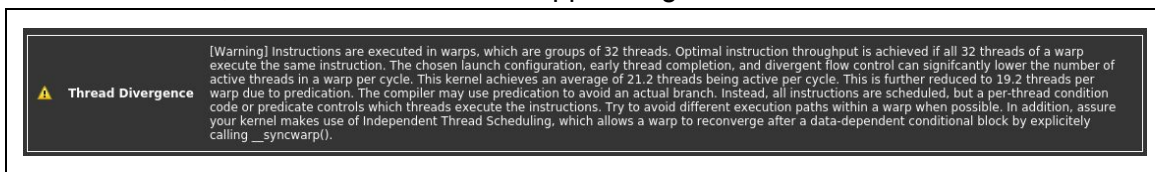




- The second kernel, ApplyBlur, which reads memory but also does comparatively many math operations, is again compute bound. Since the same pixels are read from global memory by multiple threads in a block, we see high cache hit rates for both L1 and L2. Warp stalls are dominated by MIO Throttle (math) and Short Scoreboard (Shared Memory/L1 reads).



- The third kernel, ExtractFeatures, is latency bound as indicated by the GPU Speed Of Light section warning. Inspecting the Scheduler and Warp State Statistics, we see that we issue around every 1.7 cycles (instead of every cycle). The metrics and recommendation in the Warp State Statistics section highlight the high code divergence within this kernel, which is a result of the applied algorithm.



- Inspecting the kernel's primary Sampling Data (Not Issued) on the Source page, we see that warps are mostly stalled at control flow conditions during the keypoint extraction phase (in which case control flow depends on the local pixel luminance values).

#	Source	Sampling Data (All)	Sampling Data (Not Issued)	Instructions Executed
12_	// keypoint extraction:	0	0	
12_	// get number of contiguous pixels above or below threshold	0	0	
12_	// current difference direction (brighter, darker) is stored as Luminance	0	0	
12_	int numPixels = 0;	0	0	
12_	LuminanceDiff diff = DiffUnknown;	0	0	
12_	// iterate over all pixels of the 'circle'	0	0	
12_	for (int p = 0; p < 16; ++p)	0	0	
12_	{	0	0	
12_	const auto currentLuminance = GetLuminance(pImg, tx + LUX[p], ty +	224	77	1.445.506
12_	}	0	0	
12_	if (currentLuminance < luminance - threshold)	1.129	535	2.337.152
12_	{	0	0	
12_	if (diff != DiffBrighter)	146	55	764.011
12_	{	0	0	
12_	diff = DiffDarker;	0	0	
12_	++numPixels;	24	11	154.660
12_	if (numPixels == minContPixels)	41	20	792.682
12_	{	0	0	
12_	// once we found enough pixels with the same difference	0	0	
12_	// at a time, we can stop the check	0	0	
12_	break;	0	0	
12_	}	0	0	
12_	continue;	0	0	
12_	}	0	0	
12_	else if (currentLuminance > luminance + threshold)	205	48	1.781.825
12_	{	0	0	
12_	if (diff != DiffDarker)	49	41	430.080
12_	{	0	0	
12_	diff = DiffBrighter;	0	0	
12_	++numPixels;	50	10	152.500
12_	if (numPixels == minContPixels)	320	110	1.699.130
12_	{	0	0	

See if you can find any ways to improve the performance of these kernels. Performance analysis and optimization is an iterative process that only ends when you determine the work required is not worth the return on investment. This exercise introduced several tools to help you with that process as well tools to help debug CUDA code.